



Coexecutor Runtime

1	Coexecutor Runtime	3
1.1	Introduction	
1.2	Validation	
1.3	Use cases - Examples	

1. Coexecutor Runtime

1.1 Introduction

The *Coexecutor Runtime* will be released once the research work done during the last year is published. In addition, we intend to publish several works related to the runtime, currently in progress, due to the importance and popularity that oneAPI has nowadays. We are currently studying the behavior of oneAPI for other architectures (apart from CPU and integrated GPU), as well as elaborating a new load balancing algorithm more appropriate for co-execution with this runtime, all thanks to the study of the behavior of the *Coexecutor Runtime* during its development and experimentation.

Nevertheless, if any reviewer would like to see the internal code of the *Coexecutor Runtime* to compile and execute some benchmarks, please, contact with Raúl Nozal. The code will only be provided for reviewing purposes, never to be published until we perform the official open-source releasing in the next months.

1.2 Validation

Coexecutor Runtime has been validated in another two architectures. One is a Desktop system, as a low profile computing node. Another is a powerful server node, part of a cluster infrastructure (Intel DevCloud).

```

1 CPU:
2
3     Architecture:                x86_64
4     CPU op-mode(s):              32-bit, 64-bit
5     Byte Order:                  Little Endian
6     Address sizes:                39 bits physical, 48 bits virtual
7     CPU(s):                      4
8     On-line CPU(s) list:         0-3
9     Thread(s) per core:          2
10    Core(s) per socket:          2
11    Socket(s):                   1
12    NUMA node(s):                1
13    Vendor ID:                   GenuineIntel
14    CPU family:                   6
15    Model:                       78
16    Model name:                  Intel(R) Core(TM) i5-6200U CPU @
    ↪ 2.30GHz
17    Stepping:                    3
18    CPU MHz:                     600.053
19    CPU max MHz:                 2800,0000
20    CPU min MHz:                 400,0000
21    BogomIPS:                    4801.00
22    Virtualization:              VT-x
23    L1d cache:                   64 KiB
24    L1i cache:                   64 KiB
25    L2 cache:                    512 KiB
26    L3 cache:                    3 MiB
27    NUMA node0 CPU(s):           0-3
28
29 Memory (MiB):
30          total      used      free      shared  buff/cache  available
31    Mem:    7844      927      5473        307       1443       6383
32    Swap:      0         0         0

```

Listing 1: Desktop node

```

1 CPU:
2
3   Architecture:                x86_64
4   CPU op-mode(s):              32-bit, 64-bit
5   Byte Order:                  Little Endian
6   Address sizes:               39 bits physical, 48 bits virtual
7   CPU(s):                      12
8   On-line CPU(s) list:         0-11
9   Thread(s) per core:          2
10  Core(s) per socket:          6
11  Socket(s):                   1
12  NUMA node(s):                1
13  Vendor ID:                   GenuineIntel
14  CPU family:                   6
15  Model:                       158
16  Model name:                   Intel(R) Xeon(R) E-2176G CPU @ 3.70GHz
17  Stepping:                     10
18  CPU MHz:                      4177.750
19  CPU max MHz:                  4700.0000
20  CPU min MHz:                  800.0000
21  BogomIPS:                     7399.70
22  Virtualization:              VT-x
23  L1d cache:                    192 KiB
24  L1i cache:                    192 KiB
25  L2 cache:                     1.5 MiB
26  L3 cache:                     12 MiB
27  NUMA node0 CPU(s):           0-11
28
29 Memory (MiB):
30      total      used      free      shared  buff/cache   available
31 Mem:    64133      368    63538          1        226      63179
32 Swap:   1951         0     1951

```

Listing 2: DevCloud node

1.2.1 Performance

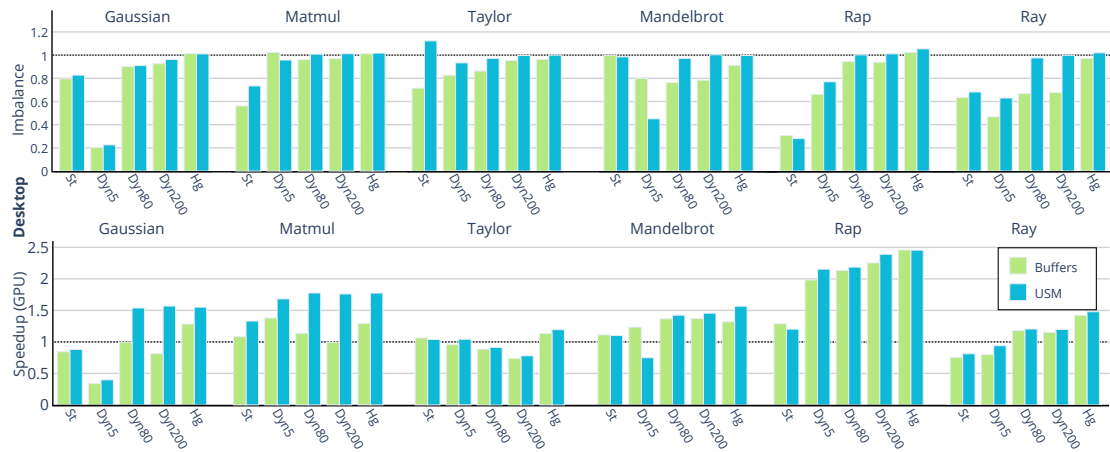


Figure 1.1: Balancing efficiency and Speedup (GPU/Coexecuting) for a set of benchmarks in a Desktop node.




Figure 1.2: Speedup (GPU/Coexecuting) for a set of benchmarks in a Devcloud server.

1.2.2 Energy consumption and efficiency

The DevCloud does not allow using neither RAPL hardware counters nor perf-utils. Therefore, its energy measurements cannot be collected and analyzed.

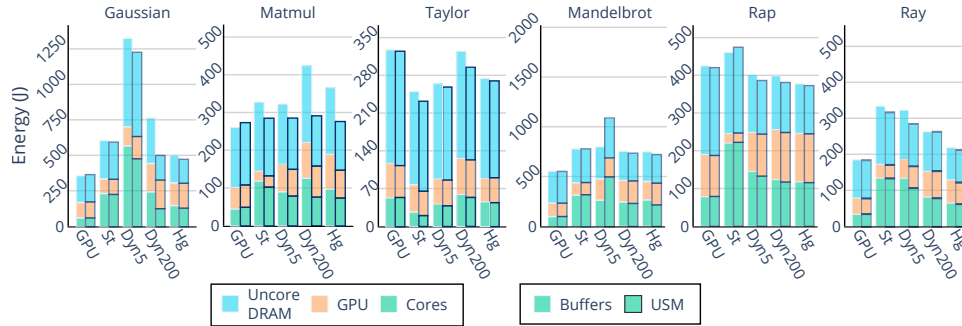


Figure 1.3: Energy consumption for a set of benchmarks in a Desktop system.




Figure 1.4: Energy efficiency (GPU/Coexecuting) for a set of benchmarks in a Desktop system.

1.2.3 Scalability

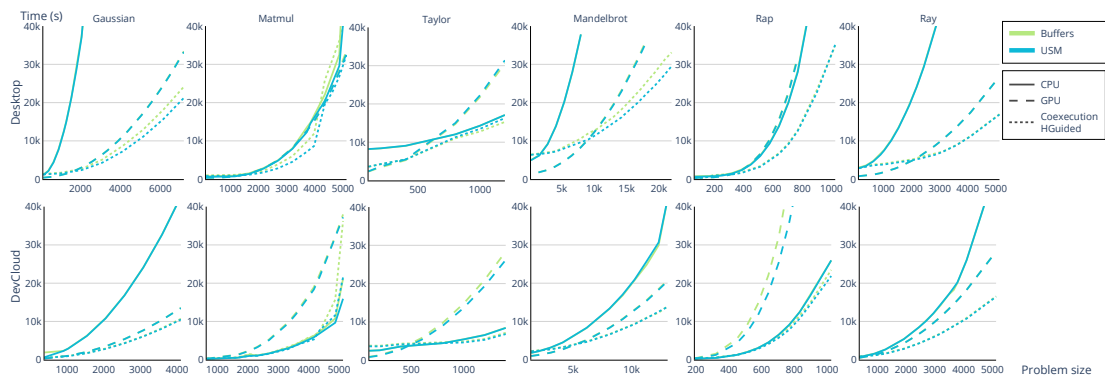


Figure 1.5: Scalability when using CPU-only, GPU-only and Coexecutor Runtime, for USM and Buffers memory models.

1.3 Use cases - Examples

1.3.1 SAXPY

Three independent SAXPY examples are depicted, showing the expressiveness of the *Coexecutor Runtime*. Two types of kernel definition are shown. One it is based on a explicit embedded context (lambda function), while the other allows accessing extended methods and callback functions, based on the *CommanderKernel* interface. In the latter, the programmer uses the *init* method and also the *callback* for kernel computation. Moreover, the *extended_computation_usm* function shows how the shared memory is managed from the point of view of the programmer. As can be seen in the three real examples, the *Coexecutor Runtime* hides all the implementation details, easing the use of its co-execution capabilities to exploit easily any oneAPI program.

```

1 class SAXPY : public CommanderKernel {
2 public:
3     SAXPY(int *x_ptr, int *y_ptr, int *out_ptr, size_t data_len,
4           float scalar_float)
5         : m_x_ptr(x_ptr), m_y_ptr(y_ptr), m_out_ptr(out_ptr),
6           m_scalar_float(scalar_float), m_data_len(data_len) {}
7
8     // pre-setup configuration - Director scope (shared internally between
9     //   ↪ Coexecution Units)
10    void init(coexecutor_unit *cunit) {
11        cunit->add_buffer<int, 1>(0, m_x_ptr, sycl::range<1>(m_data_len));
12    }
13
14    // /* @callback */
15    // void init_completed() {}
16
17    program_size size() { return m_data_len; }
18
19    void compute(coexecutor_unit *cunit, package pkg) {
20        std::cout << "[" << cunit->id() << "]" computing...\n";
21
22        sycl::buffer<int, 1> buf_x =
23            *cunit->get_buffer<int, 1>(0); // use all buffer
24        sycl::buffer<int, 1> buf_y(m_y_ptr + pkg.offset,
25            ↪ sycl::range<1>(pkg.size));
26        sycl::buffer<int, 1> buf_out(m_out_ptr + pkg.offset,
27            ↪ sycl::range<1>(pkg.size));
28
29        // communicate with the Commander to dispatch a transaction
30        // Director-Coexecution Unit communication
31        cunit->dispatch([&](sycl::handler &h) {
32            auto R = sycl::range<1>(pkg.size);
33
34            auto x = buf_x.get_access<sycl::access::mode::read>(h);
35            auto y = buf_y.get_access<sycl::access::mode::read>(h);
36            auto out = buf_out.get_access<sycl::access::mode::discard_write>(h);

```



```

36     auto scalar = (int)m_scalar_float; // caching
37
38     // Parallel region for the workload package
39     h.parallel_for(R, [=](sycl::item<1> it) {
40         auto tid = it.get_linear_id();
41         out[tid] = x[tid] * scalar + y[tid];
42     });
43 });
44 }
45
46 /* @callback */
47 void compute_completed(coexecutor_unit *counit, package pkg) {
48     // Statistics per package - related with the Coexecution Unit
49     std::cout << "[" << counit->id() << "] package " << pkg.id
50         << " computed with throughput " << pkg.throughput << "\n";
51     counit->dump_statistics();
52 }
53
54 private: // organized data:
55     int *m_x_ptr;
56     int *m_y_ptr;
57     int *m_out_ptr;
58     size_t m_data_len;
59     float m_scalar_float;
60 };
61
62 void extended_computation_usm() {
63     coexecutor_runtime<dyn> runtime;
64     runtime.config(CounitSet::CpuGpu, 12); // dynamic; 12 packages
65
66     // input data
67     size_t N = 10000000;
68     std::vector<int> x;
69     int* y;
70     std::vector<int> out(N);
71     float scalar = 3.14;
72
73     // assigns
74     x.assign(N, 1);
75     y = runtime.alloc<int>(N);
76     for (int i = 0; i < N; ++i) {
77         x[i] = x[i] + i;
78         y[i] = y[i] + 2 * i;
79     }
80     std::cout << "data x: ";
81     for (int v : x) {
82         std::cout << v << " ";
83     }
84     std::cout << "\n";

```

```

85  std::cout << "data y: ";
86  for (int i=0; i<N; ++i) {
87      std::cout << y[i] << " ";
88  }
89  std::cout << "\n";
90
91  SAXPY program{x.data(), y, out.data(), N, scalar};
92  runtime.run(program);
93
94  std::cout << "data output: ";
95  for (int v : out) {
96      std::cout << v << " ";
97  }
98  std::cout << "\n";
99
100 runtime.free<int>(y);
101 }
102
103 void extended_computation() {
104     // input data
105     size_t N = 100000000;
106     std::vector<int> x;
107     std::vector<int> y;
108     std::vector<int> out(N);
109     float scalar = 3.14;
110
111     // assigns
112     x.assign(N, 1);
113     y.assign(N, 2);
114     for (int i = 0; i < N; ++i) {
115         x[i] = x[i] + i;
116         y[i] = y[i] + 2 * i;
117     }
118     std::cout << "data x: ";
119     for (int v : x) {
120         std::cout << v << " ";
121     }
122     std::cout << "\n";
123     std::cout << "data y: ";
124     for (int v : y) {
125         std::cout << v << " ";
126     }
127     std::cout << "\n";
128
129     coexecutor_runtime<dyn> runtime;
130     runtime.config(CounitSet::CpuGpu, 20); // 20 packages given dynamically
131     SAXPY program{x.data(), y.data(), out.data(), N, scalar};
132     runtime.run(program);
133

```

```

134     std::cout << "data output: ";
135     for (int v : out) {
136         std::cout << v << " ";
137     }
138     std::cout << "\n";
139 }
140
141 void immediate_computation() {
142     // input data
143     std::vector<int> data;
144     float datav = 3.14;
145     size_t N = 200000;
146
147     // assigns
148     data.assign(N, 1);
149     for (int i = 0; i < N; ++i) {
150         data[i] = data[i] + i;
151     }
152     std::cout << "data: ";
153     for (int x : data) {
154         std::cout << x << " ";
155     }
156     std::cout << "\n";
157
158     {
159         std::mutex mut;
160         coexecutor_runtime<hg> runtime;
161         runtime.config(CountSet::CpuGpu, 0.35); // hguided; CPU has the 35%
162         ↪ of the total computation power
163         runtime.launch(data.size(), [&](coexecutor_unit *counit, package pkg)
164         ↪ {
165             {
166                 std::lock_guard<std::mutex> lk(mut);
167                 std::cout << "[" << counit->id() << "] computing " << pkg.size <<
168                 ↪ " with offset " << pkg.offset << "\n";
169             }
170
171             sycl::buffer<int, 1> buf_input(data.data() + pkg.offset,
172                 sycl::range<1>(pkg.size));
173
174             counit->dispatch([&](sycl::handler &h) {
175                 auto R = sycl::range<1>(pkg.size);
176
177                 auto input =
178                 ↪ buf_input.get_access<sycl::access::mode::read_write>(h);
179
180                 h.parallel_for(R, [=](sycl::item<1> it) {
181                     auto tid = it.get_linear_id();
182                     input[tid] = input[tid] * datav;
183                 });
184             });
185         });
186     }
187 }

```

```

179     });
180     });
181     });
182 }
183
184 std::cout << "data output: ";
185 for (int x : data) {
186     std::cout << x << " ";
187 }
188 std::cout << "\n";
189 }
190
191 int main(int argc, char *argv[]) {
192
193     immediate_computation();
194     extended_computation();
195     extended_computation_usm();
196
197     return 0;
198 }

```

1.3.2 Gaussian blur

Gaussian blur performs a gaussian filter to an image. Only the relevant lines for the programmer are shown, discarding the image and filter initialization (gaussian_cpu code). Therefore, initialization of the gaussian program and its validations are omitted for clarifying purposes. The *Coexecutor Runtime* acts transparently for the programmer, managing all the needed resources regarding oneAPI primitives and scheduling behavior.

```

1  #include "gaussian.h"
2
3  int main(int argc, char *argv[]) {
4      // setup initial data, prepare filter and image
5      int filterSize = 61;
6      size_t N = 2048;
7      int imageWidth = N, imageHeight = N;
8      gaussian_cpu gaussian(imageWidth, imageHeight, filterSize);
9
10     uchar4 *a_ptr = (uchar4 *)malloc(gaussian._total_size * sizeof(uchar4));
11     float *b_ptr = (float *)malloc(gaussian._filter_total_size *
12     ↪     sizeof(float));
13     uchar4 *c_ptr = (uchar4 *)malloc(gaussian._total_size * sizeof(uchar4));
14
15     gaussian.set_buffers(a_ptr, b_ptr, c_ptr);
16     gaussian.build(); // compose images
17
18     {
19         coexecutor_runtime<dyn> runtime; // Coexecutor-Runtime
20         // use the dynamic scheduler with CPU-GPU

```

```

20 // co-execution using 32 packages
21 runtime.config(CounitSet::CpuGpu, 32);
22 runtime.launch(
23     gaussian._total_size,
24     [&](coexecutor_unit *counit, package pkg) { // launch the Gaussian
25         ↪ bench
26         std::cout << "[" << counit->id() << "] computing " << pkg.size
27             << " with offset " << pkg.offset << "\n";
28
29         auto Rinput = sycl::range<1>(gaussian._total_size);
30         auto Rfilter = sycl::range<1>(gaussian._filter_total_size);
31         auto cols = imageWidth;
32         auto rows = imageHeight;
33
34         sycl::buffer<uchar4, 1> buf_input(gaussian._a, Rinput);
35         sycl::buffer<float, 1> buf_filterWeight(gaussian._b, Rfilter);
36         sycl::buffer<uchar4, 1> buf_blurred((gaussian._c + pkg.offset),
37             sycl::range<1>(pkg.size));
38
39         counit->dispatch([&](sycl::handler &h) {
40             auto R = sycl::range<1>(pkg.size);
41
42             auto input =
43                 ↪ buf_input.get_access<sycl::access::mode::read>(h);
44             auto filterWeight =
45                 ↪ buf_filterWeight.get_access<sycl::access::mode::read>(h);
46             auto blurred =
47                 ↪ buf_blurred.get_access<sycl::access::mode::discard_write>(h);
48
49             h.parallel_for(R, [=](sycl::item<1> it) {
50                 auto tid = it.get_linear_id() + pkg.offset;
51                 int r = tid / cols;
52                 int c = tid % cols;
53
54                 int middle = filterSize / 2;
55                 float blurX = 0.f;
56                 float blurY = 0.f;
57                 float blurZ = 0.f;
58
59                 int width = cols - 1;
60                 int height = rows - 1;
61
62                 for (int i = -middle; i <= middle; ++i) {
63                     for (int j = -middle; j <= middle; ++j) {
64
65                         int h = r + i;
66                         int w = c + j;
67                         if (h > height || h < 0 || w > width || w < 0) {

```

```

66         continue;
67     }
68
69     int idx = w + cols * h;
70
71     float pixelX = (input[idx].x());
72     float pixelY = (input[idx].y());
73     float pixelZ = (input[idx].z());
74
75     idx = (i + middle) * filterSize + j + middle;
76     float weight = filterWeight[idx];
77
78     blurX += pixelX * weight;
79     blurY += pixelY * weight;
80     blurZ += pixelZ * weight;
81 }
82 }
83
84     tid -= pkg.offset;
85     blurred[tid].x() = (unsigned char)cl::sycl::round(blurX);
86     blurred[tid].y() = (unsigned char)cl::sycl::round(blurY);
87     blurred[tid].z() = (unsigned char)cl::sycl::round(blurZ);
88 });
89 });
90 });
91
92     // show runtime info
93     runtime.dump_statistics();
94 }
95
96 if (gaussian.compare_gaussian_blur(THRESHOLD)) { // unit test
97     std::cout << "Success\n";
98 } else {
99     std::cout << "Failure\n";
100 }
101
102 // explicit free memory
103 free(a_ptr);
104 free(b_ptr);
105 free(c_ptr);
106
107 return 0;
108 }

```
